

```
In [0]: 1 from google.colab import drive
        2 drive.mount('/content/drive/', 'My Drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code

Enter your authorization code:
.....

Mounted at /content/drive/

▼ 1 Data downloaded from kaggle <https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge> (<https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge>)

```
In [0]: 1 %matplotlib inline
        2 %matplotlib notebook
        3 import pandas as pd
        4 import numpy as np
        5 import matplotlib.pyplot as plt
        6 from sklearn.feature_extraction.text import CountVectorizer
        7 from sklearn.feature_extraction.text import TfidfVectorizer
        8 from sklearn.model_selection import train_test_split
        9 from scipy.sparse import hstack
       10 from sklearn.linear_model import LogisticRegression
       11 from sklearn.model_selection import cross_val_score
       12 from sklearn.metrics import roc_auc_score
       13
```

```
In [0]: 1 data = pd.read_csv('/content/drive/My Drive/ToxicComments/data/train.csv')
```



In [0]: 1 data.head()

Out[4]:

	id	comment_text	toxic	severe_toxic	obscene	threat	insult	iden
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	
1	000103f0d9cfb60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	

In [0]: 1 data.drop(columns=['id'], axis=1, inplace=True)

In [0]:

```

1 def objectionableOrNot(dataRow):
2     if (dataRow['toxic'] == 1 or dataRow['severe_toxic'] == 1 or dataRow['obscene'] == 1 or
3         dataRow['insult'] == 1 or dataRow['identity_hate'] == 1) :
4         retVal = 1
5     else:
6         retVal = 0
7     return retVal

```

In [0]: 1 list(data.columns[1:7])

Out[7]: ['toxic', 'severe_toxic', 'obscene', 'threat', 'insult', 'identity_hate']

```
In [0]: 1 data['objectionable'] = data.apply(objectionableOrNot, axis=1)
2 data.drop(columns=list(data.columns[1:7]), axis=1, inplace=True)
3 data.head()
```

```
Out[6]:
```

	comment_text	objectionable
0	Explanation\nWhy the edits made under my usern...	0
1	D'aww! He matches this background colour I'm s...	0
2	Hey man, I'm really not trying to edit war. It...	0
3	"\nMore\nI can't make any real suggestions on ...	0
4	You, sir, are my hero. Any chance you remember...	0

```
In [0]: 1 data.objectionable = data.objectionable.astype('category')
```

```
In [0]: 1 data[data.objectionable == 1].shape
```

```
Out[48]: (16225, 2)
```

```
In [0]: 1 data[data.objectionable == 0].shape
```

```
Out[49]: (143346, 2)
```

```
In [0]: 1 train_x, test_x, train_y, test_y = train_test_split(data['comment_text'], d
2 print(f'Train size {train_x.shape}')
3 print(f'Train size {test_x.shape}')
```

```
Train size (111699,)
Train size (47872,)
```

```
In [0]: 1 # upsampling
2 # Indices of each class' observations
3 i_class0 = np.array(train_y[train_y == 0].index)
4 i_class1 = np.array(train_y[train_y == 1].index)
5
6 # Number of observations in each class
7 n_class0 = len(i_class0)
8 n_class1 = len(i_class1)
9
10 # For every observation in class 0, randomly sample from class 1 with replac
11 i_class1_upsampled = np.random.choice(i_class1, size=n_class0, replace=True)
```

```
In [0]: 1 print(n_class0)
2 print(n_class1)
```

```
100342
11357
```

```
In [0]: 1 train_x = pd.concat([train_x.loc[i_class1_upsampled], train_x.loc[i_class0]
```

```
In [0]: 1 train_y = pd.concat([train_y.loc[i_class1_upsampled], train_y.loc[i_class0]
```

```
In [0]: 1 tfidf_vectorizer_word = TfidfVectorizer(strip_accents='unicode', analyzer='
2 train_x_tfidf_word = tfidf_vectorizer_word.fit_transform(train_x)
3 test_x_tfidf_word = tfidf_vectorizer_word.transform(test_x)
```

```
In [0]: 1 type(train_x_tfidf_word)
```

```
Out[15]: scipy.sparse.csr.csr_matrix
```

```
In [0]: 1 from scipy.sparse.csr import csr_matrix
```

```
In [0]: 1 tfidf_vectorizer_char = TfidfVectorizer(strip_accents='unicode', stop_words
2 train_x_tfidf_char = tfidf_vectorizer_char.fit_transform(train_x)
3 test_x_tfidf_char = tfidf_vectorizer_char.transform(test_x)
```

```
In [0]: 1 train_tfidf_word_char = hstack([train_x_tfidf_word, train_x_tfidf_char])
2 test_tfidf_word_char = hstack([test_x_tfidf_word, test_x_tfidf_char])
```

```
In [0]: 1 train_tfidf_word_char = train_tfidf_word_char.tocsr()
2 test_tfidf_word_char = test_tfidf_word_char.tocsr()
```

```
In [0]: 1 from sklearn.naive_bayes import MultinomialNB
```

```
In [0]: 1 def predict(review, clf):
2     char_feature = tfidf_vectorizer_char.transform([review])
3     word_feature = tfidf_vectorizer_word.transform([review])
4     features = hstack([char_feature, word_feature]).tocsr()
5     return clf.predict(features)
```

```
In [0]: 1 mnbn = MultinomialNB(alpha=0.1)
2 mnbn.fit(train_tfidf_word_char, train_y)
```

```
Out[91]: MultinomialNB(alpha=0.1, class_prior=None, fit_prior=True)
```

```
In [0]: 1 print(list(np.where(test_y ==1)[0:20]))
```

```
[array([ 22,  45,  51, ..., 47856, 47866, 47867])]
```

```
In [0]: 1 y_pred_proba = mnbn.predict_proba(test_tfidf_word_char)
2 roc_auc_score(test_y, y_pred_proba[0:, 1], average='micro')
```

```
Out[92]: 0.9476282928946549
```

```
In [0]: 1 ▾ for r in np.where(test_y ==1)[0][0:20]:  
        2     print(predict(test_x.iloc[r], mnb))
```

```
[0]  
[1]  
[1]  
[1]  
[1]  
[1]  
[0]  
[1]  
[1]  
[1]  
[1]  
[0]  
[1]  
[1]  
[1]  
[1]  
[1]  
[1]  
[1]  
[1]  
[1]
```

```
In [0]: 1 import pickle
```

```
In [0]: 1 pickle.dump(tfidf_vectorizer_word, open('/content/drive/My Drive/ToxicCommentsClassification/tfidf_vectorizer_word.pkl', 'wb'))  
        2 pickle.dump(tfidf_vectorizer_char, open('/content/drive/My Drive/ToxicCommentsClassification/tfidf_vectorizer_char.pkl', 'wb'))  
        3 pickle.dump(mnb, open('/content/drive/My Drive/ToxicCommentsClassification/mnb.pkl', 'wb'))
```